

# Data Quality and Observability System for Financial Data Pipelines

*Ensuring data integrity, correctness, and reliability is critical for a personal finance app that handles both personal and business transactions. Below is a comprehensive system covering validation rules, observability practices, data lineage, and alerting policies to catch errors, drifts, and silent failures in financial intelligence pipelines.*

## 1. Validation Rules (Data Quality Checks)

A robust **validation rule catalog** should enforce business logic and data integrity at every step of the pipeline. Key validation rules include:

- **Balance Integrity (No Negative Net):** Verify that financial balances and net sums never drop below zero (except where logically allowed, e.g. debt accounts). This acts as a range check to catch impossible values (e.g. a net worth or account balance that is negative when it shouldn't be). Range checking ensures numeric values stay within acceptable boundaries and is *"particularly valuable for financial figures and any metric with natural limits"* <sup>1</sup>. For example, if an account's closing balance or a forecasted cash flow goes negative without a corresponding liability, the system should flag it for review.
- **Referential Integrity (No Orphans):** Ensure every record that should have a parent or reference (such as a payout or transfer) is linked to a valid transaction or account. No "orphan" payouts should exist without an originating transaction reference. Enforcing referential integrity catches any orphaned records and broken links in the data. For instance, if there's a payout logged that doesn't tie back to any expense/income entry, it should be flagged. This type of check *"validates connections between related data... Every order must link to an existing customer... These checks prevent orphan records that clutter databases"* <sup>2</sup> (in our case, every payout or invoice entry must link to a transaction or account).
- **Duplicate & Loop Transaction Elimination:** Detect and eliminate duplicate entries (or circular references) to avoid double-counting. Duplicate transactions – for example, the same bank transaction imported twice or a transfer recorded as both an expense and income – can skew projections and analytics. A uniqueness check should enforce that transactions meant to be unique (by ID or timestamp+amount combination) are not repeated <sup>3</sup>. Additionally, handle internal transfers carefully: if money moving between personal and business accounts is recorded in both, the system should recognize it as a transfer (netting out) rather than two separate income/expense events, preventing a "loop" in which a transfer inflates both income and expense. Automated de-duplication rules (by transaction ID or fuzzy matching of amount/date) and business logic to identify transfers ensure we don't count the same money twice.

- **Required Fields and Valid Formats:** Validate that all essential fields are present and correctly formatted. Every transaction should have a date, amount, category, and account specified. Missing critical data (null or blank fields) should be flagged immediately <sup>4</sup>. For example, a transaction with no category or a null amount is invalid for analysis. Type and format checks should enforce that dates are valid (no February 30th, etc.), amounts are numerical, and text fields meet expected patterns. These basic schema checks prevent downstream issues (e.g., “N/A” or an out-of-range date causing errors) <sup>5</sup> <sup>6</sup>.
- **Category Validity & Consistency:** Since categorization is crucial for separating personal vs business finances, enforce that each transaction’s category is one of the allowed values (personal expense, business expense, income, etc.) and aligns with the transaction type. If a transaction is marked as a business expense but belongs to a personal account, or if an AI-generated category is “unknown”, the system should flag it. This can be implemented with *accepted value* rules (e.g. category must belong to the predefined category list) similar to how one would enforce an age field to be non-negative <sup>7</sup> <sup>8</sup>. Consistency checks can also cross-verify fields – for example, if a transaction is in a personal account, its category should be a personal category (not a business category), ensuring data stays consistent with the user’s personal/business separation logic.
- **Balancing & Aggregation Rules:** In a double-entry or aggregated setting, implement checks to ensure sums align. For instance, in business accounting features, ensure debits equal credits if applicable (no unbalanced journal entries) <sup>9</sup>. For budgets or projections, verify that total incomes minus expenses equals net savings for the period, without unexplained discrepancies. Any breach (like expenses exceeding income by an unrealistic amount without drawing from savings or credit) should trigger investigation. These higher-level business rules act as “sanity checks” on the outputs of the pipeline.

All the above rules can be automated using open-source data validation frameworks. For example, **Great Expectations** (an open-source tool) allows you to declaratively define these expectations (no negatives, no nulls, unique transaction IDs, etc.) and will flag or even stop the pipeline if they fail. **dbt** (data build tool) also supports tests for non-null, uniqueness, accepted values, and referential integrity out-of-the-box <sup>10</sup> <sup>8</sup>. Using such frameworks, we can maintain a **Validation Rule Catalog** documenting each rule, its purpose, and how it’s implemented (SQL tests, Python checks, etc.), ensuring every data ingestion or transformation step is guarded by quality gates. This catalog becomes the reference for data quality in the system, and each rule is tied to an alert (discussed below) so issues are immediately visible.

## 2. Observability and Monitoring

To catch silent failures and drifts early, we need **observability** across both the data pipelines and the data itself. This involves comprehensive monitoring of pipeline health, data quality metrics, and ML model performance. Key aspects include:

### Pipeline & Sync Job Health Dashboards

Monitor the **operational health** of all ETL/ELT jobs that sync and process data. Every scheduled job (e.g. nightly bank transaction sync, categorization runs, forecasting jobs) should report metrics to a centralized dashboard. At minimum, track: **success/failure status**, run durations, data volumes processed, and timeliness of execution. Observability is not just “did the job fail?”, but also “*Did it run on time? Did it process*

the expected amount of data? Did it create correct outputs? Are data volumes within thresholds?" <sup>11</sup> . For example, if a daily import job usually processes ~100 transactions but today processed 0 or 1000, that anomaly should be evident on the dashboard.

A good practice is to implement three layers of monitoring <sup>12</sup> :

- **Operational monitoring:** Track workflow status and performance – e.g. each pipeline's run completion (success/failure), start/end timestamps, runtime, any SLA misses (was it delayed or faster than expected) <sup>13</sup> . Tools like Apache **Airflow** (open source) provide a built-in UI to see DAG (Directed Acyclic Graph) runs, task failures, and retry counts, and can be configured with alerts on failures or SLA misses <sup>14</sup> . This ensures any job failure or slowdown (perhaps due to an upstream API issue or code bug) is immediately visible.
- **Data monitoring:** Even if a job succeeds, ensure the data output is within expected bounds. This means tracking row counts, distribution of values, schema changes, etc. For instance, if the "transactions" table suddenly has 20% fewer rows than yesterday or a schema column like "amount" changed type, it could indicate a partial load or upstream change. **Null or empty field rates, spikes/drops in transaction counts, out-of-range values, and duplicate counts** are all data metrics to watch <sup>15</sup> . A monitoring dashboard might show daily total transactions, percentage of transactions categorized as "uncategorized", etc., with thresholds for normal ranges.
- **System monitoring:** Watch the infrastructure (memory, CPU, etc.) if you self-host pipeline components, to catch resource issues that could impact data jobs <sup>16</sup> . For example, if using a Python script on a server to ingest data, a memory leak could eventually cause silent truncation of data – system metrics help correlate such issues.

**Open-source observability tools** can be leveraged here. A common stack is **Prometheus** for metrics collection and **Grafana** for visualization and alerting <sup>17</sup> <sup>18</sup> . Instrument your pipeline jobs to push metrics (e.g. count of records processed, job duration, last run time, number of errors) to Prometheus. Then create Grafana dashboards to display these metrics in real time. Grafana can show, for example, a **"Sync Job Health"** dashboard with graphs for each job's runtime and rows processed each day, and use color-coding to highlight anomalies. This way, at a glance, you can see if last night's data sync behaved normally. As one guide suggests, *"Use Prometheus to collect metrics from your data pipelines and systems. Configure Grafana to visualize these metrics and set up alerts for critical thresholds."* <sup>19</sup> . For instance, if a pipeline hasn't run by its scheduled time or processed zero records, Grafana can raise an alert (through Slack/email as detailed later).

## Forecast Accuracy Monitoring

For any predictive or forecasting components (e.g. projecting future expenses, cash flow, or budget forecasts), implement continuous **model performance monitoring**. It's not enough to generate forecasts – we must verify they stay accurate over time. As actual data comes in (e.g. the actual expenses for the month), compare it against the forecast and quantify the error. Key metrics like Mean Absolute Percentage Error (MAPE) or Mean Absolute Error can be computed for each forecast cycle. These accuracy metrics should be tracked on a dashboard to spot when predictions start diverging from reality.

For example, if the app forecasts that next month's spending on groceries will be \$500, but the actual turns out to be \$800, that's a large error. Monitoring the trend of forecast errors across months can reveal if the

forecasting model is degrading. If you notice the error metrics creeping above a threshold (say, MAPE > 20% consistently), that should trigger retraining the model or investigating the cause (maybe an economic change or one-off event).

This kind of monitoring falls under **ML model observability**, which focuses on how the model behaves in production beyond just technical uptime <sup>20</sup>. It's recommended to track *"metrics related to model quality (e.g., accuracy, precision, etc.)"* as part of observability <sup>21</sup>. In our context, that means tracking the accuracy of categorization (if the model suggests categories) and the accuracy of any financial forecasts. The monitoring system can log each prediction made by the model and later log the actual outcome once known, computing error metrics automatically. These metrics can be sent to the same Prometheus/Grafana setup or specialized ML monitoring tools (there are open-source options like **Evidently** for model monitoring). The goal is to have a **Forecast Performance Dashboard** that might show, for instance, a line chart of forecast vs actual for key metrics (monthly expense total, cash flow, etc.) and a rolling accuracy percentage. If accuracy dips below an acceptable threshold, it should raise an alert (preventing "silent" model decay where the model's suggestions quietly become unreliable).

## Drift Detection in ML Categorizations

If the system uses an AI/ML model to categorize transactions or an "AI agent" to make financial recommendations, it's vital to monitor for **data drift** and **concept drift** in those models. Data drift refers to changes in the input data's distribution over time <sup>22</sup>. For example, perhaps the user's spending patterns change (new merchants, different frequencies) or the model was trained on last year's data which isn't representative of this year. **Drift detection** mechanisms should be in place to catch these shifts before they cause major misclassifications.

One practical approach is to log the distribution of predicted categories (and possibly key features like transaction amounts or merchant types) and compare the current distribution to a baseline (such as the training data distribution or last quarter's distribution). If the model suddenly starts classifying 40% of transactions as "Other" whereas previously it was 5%, that's a red flag. Statistically, we can use tests like **chi-square for categorical data** or **Kolmogorov-Smirnov for numeric features** to detect if the new data distribution is significantly different from the old <sup>23</sup>. In fact, data observability frameworks often include drift detection that *"compares the distribution of new incoming data to the original training data using statistical tests"* <sup>24</sup>. Open-source tools like **Evidently** provide ready-made monitors for data drift; Evidently's library can compute a "drift score" and even generate reports showing which features or categories have shifted. According to Evidently's documentation, *"you can use tests like Kolmogorov-Smirnov for numerical features or the Chi-square test for categorical attributes"* to assess drift <sup>23</sup>.

In practice, for our financial app, we would: (a) establish a baseline distribution of transaction categories (and other features the ML model uses) from a stable period, (b) each week or month, compare the latest data against the baseline, and (c) trigger an alert if drift is detected beyond a certain confidence level. For example, if p-value of the chi-square test for category distribution is below 0.01, we suspect drift. The observability dashboard can include a **"Model Drift Monitor"** panel that perhaps shows a drift score or a comparison of category frequency histograms over time, so we can visually see changes. This helps ensure the categorization model remains accurate; if the input data characteristics change (e.g. new types of transactions post-pandemic, or a user's business starts dealing with a new expense category), we'll catch it and can retrain or adjust the model accordingly. By monitoring drift, we essentially put a safety net under

the ML aspects of the pipeline so that *“the model still performs as expected and operates under familiar conditions”* <sup>22</sup> – if not, we’ll know quickly.

## End-to-End Observability Stack

To tie it together, it’s ideal to integrate pipeline job monitoring with data quality checks and ML monitoring into one observability stack. A possible open-source stack might be:

- **Data pipeline orchestrator + tests:** Use something like **Apache Airflow** or **Dagster** to schedule and run pipeline tasks. These tools have built-in observability features: Airflow’s UI tracks task runs and can trigger callbacks on failures, and Dagster provides asset materialization tracking and even integrates with data quality checks like Great Expectations and can capture lineage <sup>25</sup>. Both allow logging and metrics export. For instance, Airflow can be configured to send metrics to StatsD/Prometheus and has operators for Slack alerts on failures <sup>14</sup>.
- **Great Expectations or Soda:** Include data validation steps in the pipeline that output reports if expectations fail. Great Expectations can be set to output a JSON or HTML report of validation results each run. These results can feed into the observability system (e.g., count of failed expectations can be a metric). Soda Core is another open-source tool that can run SQL checks and anomaly detection on data and send alerts <sup>26</sup>.
- **Metrics & Dashboarding:** As mentioned, **Prometheus** for scraping metrics (from jobs, from custom scripts, from Great Expectations results) and **Grafana** for dashboards. Grafana will display real-time views and can also serve as the central alerting mechanism (via its Alertmanager or integrating with something like Grafana OnCall or PagerDuty for incidents).
- **Log aggregation:** Centralize logs from all pipeline components (maybe using an ELK stack or Grafana Loki). This way, if something goes wrong, engineers can quickly search logs. For observability, having logs, metrics, and traces in one place is ideal (**OpenTelemetry** can be used to add tracing – instrument code so that as data flows, you can trace requests through the system <sup>27</sup>, although for a smaller-scale pipeline, this might be optional).

By implementing these layers, **silent failures** become loud – e.g., if a job technically “succeeds” but processes 0 records or produces anomalous results, the dashboards and alerts will make sure it’s noticed. As one engineer put it, *“You’re not just monitoring pipelines – you’re monitoring trust.”* <sup>28</sup> The data observability setup gives confidence that any deviation in pipeline behavior or data correctness will surface quickly.

## 3. Data Lineage and Traceability

Implementing **data lineage** is crucial for transparency and trust in a financial data system. Data lineage means having the ability to trace any derived data (a dashboard metric, a forecast, a categorized transaction) all the way back to the raw inputs that produced it, and to understand each transformation in between <sup>29</sup>. In practice, this involves capturing metadata about how data flows through the pipelines.

**Traceability from Dashboard KPIs to Raw Transactions:** Every key number on a user’s dashboard (say, “Total Personal Expenses this month” or a forecasted savings figure) should be reproducible and explainable by drilling down into source data. For example, if the dashboard says \$5,000 was spent on business expenses this quarter, the system should be able to trace that to the list of transactions contributing to that sum. One way to enable this is by maintaining references or IDs through transformations – e.g., each aggregated record can carry a list of source transaction IDs or a group identifier so that you can retrieve source records. Alternatively, using a data warehouse with a well-designed schema, you might maintain materialized views for KPIs that can be joined back to transaction tables.

A **data lineage diagram** helps illustrate these relationships. Below is a Mermaid diagram showing a simplified lineage from raw data sources to final dashboard metrics:

```

graph TD
    subgraph Sources
        A[Bank Transactions API/CSV]
        B[Manual Inputs (Payouts, Invoices)]
    end
    subgraph Pipeline["ETL/Processing Jobs"]
        A --> C1["ETL1[Ingestion & Cleaning]"]
        C1 --> C2["C[(Raw Transactions Table)]"]
        B --> ETL1["ETL1"]
        C2 --> ETL2["ETL2[ML Categorization Model]"]
        ETL2 --> D["D[(Categorized Transactions)]"]
        D --> ETL3["ETL3[Forecasting Model]"]
        ETL3 --> E["E[(Forecast Results)]"]
    end
    subgraph Analytics["Analytics & Dashboards"]
        D --> KPI1["KPI1[Monthly Spending KPI]"]
        E --> KPI2["KPI2[Projected Cashflow KPI]"]
    end

```

*Figure: Example data lineage from raw inputs to analytics.* In this example, data from external **Sources** (bank feeds and user inputs) flow through a series of **Pipeline** transformations (ingestion, categorization, forecasting). Intermediate datasets like the cleaned raw transactions table, categorized transactions, and forecast results are stored (e.g., in a database or warehouse). Finally, those feed into **Analytics/Dashboards** where KPIs are displayed. With lineage tracking, if a KPI seems off, you can trace back: the “Monthly Spending KPI” is derived from the Categorized Transactions, which in turn come from specific raw transactions. Each arrow implies a relationship that can be documented.

To implement lineage in practice, consider using an open standard or tool: **OpenLineage** is an open-source standard for tracking data lineage across jobs, supported by tools like Airflow and Spark. Tools like **Marquez** (OpenLineage’s reference implementation) can collect lineage metadata (e.g., which dataset/table was produced by which job and what source it read). If using **Dagster**, it has built-in concept of assets and can automatically record how assets (datasets) are produced from others, effectively creating a lineage graph. Even without specialized tools, you can start by maintaining a simple log or table that records, for each pipeline run, what source data it read and what outputs it wrote. The aim is to have a “*map of a dataset’s path from ingestion to visualization... an essential component of Data Observability because it empowers*”

*engineers to quickly trace the root cause of a data quality incident back to its source”* <sup>30</sup> . In other words, when something goes wrong or looks wrong, lineage lets you pinpoint where in the chain the problem occurred.

**Change-Set Tracking (Audit Logs):** In financial systems, it’s important to know *who* did *what* and *when* to the data. Implement an audit logging mechanism for data changes and pipeline changes. For example, if a user or an admin updates a transaction’s category (reclassifying an expense from personal to business), log that event (old value, new value, timestamp, user). If a pipeline process modifies or backfills data, have the code log the details (which records, what changed). These audit logs provide a timeline of changes that is invaluable for debugging and compliance. They answer questions like: “When was this forecast last updated and by which process?” or “Who corrected this transaction amount?”. Storing these logs in an immutable store (append-only logs) ensures an accurate history.

On the pipeline side, use version control for code and configurations, and consider tagging data outputs with pipeline run IDs or git commit hashes. That way, if a metric on the dashboard came from a certain run of the forecasting job, you can trace exactly which code version and input data were involved. This practice aligns with the idea that *financial regulations demand audit trails for how key metrics are calculated* – auditors should be able to see the lineage and transformation of data without manual guesswork <sup>31</sup> . Having systematic lineage and audit logs *“reduces the engineering effort required for compliance tasks from days of manual investigation to minutes of lineage lookup”* <sup>32</sup> .

**Explainability for AI Agent Actions:** If your system includes an AI agent (for example, an automated assistant that categorizes transactions or makes financial recommendations autonomously), incorporate **explainability and traceability** for its decisions. Every action the AI takes should be logged in detail – input, output, and reasoning (if available). For instance, if an AI re-categorizes a transaction or adjusts a forecast, log what triggered it and the before/after. This might include logging the prompt or features the AI saw and the result it produced. Having an *“immutable audit trail of every agent action—from the original prompt to tool execution”* is critical <sup>33</sup> . This ensures that if the AI does something unexpected, you can reconstruct why it happened. In addition to logging, aim to provide user-facing explainability where possible: e.g., show the user a note like “Transaction categorized as **Travel** by AI due to description ‘Airbnb’ (confidence 90%).” Such transparency builds trust and allows manual correction if the AI is wrong.

In more advanced setups, you might integrate **explainable AI** techniques: if using a machine learning model, store feature importance or model confidence along with each prediction. For example, if an ML model predicts a transaction is a business expense, record that it was, say, 95% confidence and top factors were “Vendor == Office Depot”. This kind of metadata can be surfaced in debugging tools or even to the user in an explanation view. The goal is to never have a “black box” in the pipeline – every automatic decision can be traced and explained. This not only helps in understanding and debugging the AI agent’s behavior, but also guards against the agent silently failing. If the agent starts doing something odd (due to drift or a bug), the combination of lineage and action logs will let you pinpoint the issue and even roll back if needed. In fact, some modern solutions even allow “replaying” or undoing AI actions using logged history <sup>34</sup> <sup>35</sup> – while that might be advanced, at minimum our system should log actions well enough that we could manually revert any undesired changes the AI made.

By implementing data lineage, audit logs, and AI explainability, **traceability** is achieved at two levels: **data traceability** (for the data itself through transformations) and **decision traceability** (for automated actions and decisions). This end-to-end lineage and logging greatly aids debugging (when a dashboard number looks wrong, lineage is your map to find where the issue arose <sup>36</sup> ) and builds confidence for users and

regulators. Business users and auditors can be given read-access to lineage information to answer questions like “Why is this number on my report, and can you prove its accuracy?” – and the system will have the answer backed by data.

## 4. Alerting and Notifications

All the validation checks and observability metrics in the world are only useful if someone knows when something is wrong. That’s where a robust **alerting system** comes in. We define clear thresholds for anomalies, and when they are breached, automated notifications are dispatched via Slack, email, or other channels, with an **escalation policy** to ensure resolution.

**Threshold-Based Alert Triggers:** Each kind of data quality rule or observability metric should have associated alert criteria. Some examples of alert triggers in this financial pipeline context:

- **Data Validation Failures:** If any of the validation rules from Section 1 fail beyond an acceptable threshold, trigger an alert. For instance, if *Balance Integrity* check fails (a negative balance is detected) even once, that could be critical (since it might indicate a serious error in calculations – perhaps double-counting debt). Or if duplicate transactions are detected above a small tolerance (say more than 1 duplicate in a batch), trigger an alert to investigate data sources. These alerts help catch issues like data corruption or logic errors early. They can be implemented by integrating the validation framework with alerting – e.g., Great Expectations can be run as part of the pipeline and if an expectation fails, have a step that sends an alert (via Slack API or email). In Airflow, one might use a failure callback to send a Slack message if a data quality check task fails.
- **Pipeline Failures and Delays:** Any ETL job failure should immediately page or notify the team. For example, if the nightly sync job fails to run or errors out, send a **high-priority** alert (Slack message to #data-team channel and email to the engineer on-call). Similarly, if a job is running but exceeds its normal runtime significantly (suggesting it’s stuck or slowed), you might alert. Airflow and other orchestrators support SLA miss alerts – e.g., Airflow can be configured to trigger an alert if a job doesn’t complete by a certain time <sup>37</sup>. These ensure “silent” pipeline issues (like a job hung waiting on an external API) don’t go unnoticed.
- **Data Drift or Anomaly Alerts:** Observability metrics like a sudden drop in number of transactions, a spike in unclassified transactions, or detected ML drift should raise alerts. For example, if drift detection (Section 2) finds a statistically significant drift in transaction categories, that could be an early warning of model issues or data changes – send an alert to the data science team to review. Another example: if forecast accuracy for the latest month falls below, say, 80%, trigger an alert that the forecast may need retraining. Data volume anomalies can be caught similarly: “*Data anomaly: row count < 10% of baseline*” can be detected by tools like Great Expectations or custom SQL, and hooked to Slack <sup>38</sup>. In fact, a guide suggests using Great Expectations with Slack for such data anomaly alerts <sup>38</sup>.
- **Infrastructure/Sync Issues:** If using external APIs or feeds (like bank APIs), an alert should fire if data hasn’t arrived in a given window (data freshness alerts). For instance, “No new transactions in last 24 hours” could indicate an upstream issue. Also, if system metrics show, say, disk nearly full or other issues that could soon impact the pipeline, those should alert the devops team.



Each alert should have a defined **severity level** (info, warning, critical) and a documented **owner** (who is responsible to respond). For a personal finance app, critical data issues might be handled by the engineering team, whereas some lower-level alerts might just log or notify an admin to check later. The motto is to avoid “alert fatigue” by tuning thresholds so that we only alert on meaningful deviations <sup>39</sup> <sup>40</sup> .

**Notification Channels (Slack/Email/Agents):** We integrate the alert triggers with communication channels to ensure rapid response. **Slack** is ideal for real-time team alerts – e.g., a Slack bot can post in a channel “*Data Pipeline Failure: Daily transactions sync failed at 2:00 AM – click here for logs.*” Use Slack webhooks or built-in operators (Airflow has `SlackWebhookOperator`, etc.) to send these. For high-severity incidents (pipeline down, data quality issue impacting many users), also send **Email** to a broader list or specific on-call individual, as email provides a persistent record. Email can also be used for external notifications (if needed to alert a business user or stakeholder that data might be delayed, though usually internal alerts suffice).

The system can also interface with an **on-call rotation or incident management** system. Open-source options like **Grafana OnCall** or commercial ones like PagerDuty can manage on-call schedules and escalation. For example, you could configure that a critical alert first sends a Slack message; if no one acknowledges in 15 minutes, Grafana OnCall escalates by emailing/texting a specific person or the next tier support <sup>41</sup> <sup>42</sup> . This ensures that urgent issues (like a broken pipeline) wake someone up if they occur off-hours. Grafana OnCall documentation describes features like “*Threshold-based escalation: escalate only if a certain number of alerts occur in a time frame*” and “*Repeat escalation: loop the alert until acknowledged*” <sup>43</sup> , which can be applied to avoid situations where an alert gets missed. In our context, an example escalation path might be: if a critical nightly job fails, send Slack alert immediately; if not acknowledged in 10 minutes, send SMS to the engineer on call; if 30 minutes pass with no acknowledgment, call the engineering manager. We tailor these paths based on the severity and business impact of each alert.

Additionally, the mention of **agent notifications** suggests we can integrate alerts with an AI agent or automated assistant if one exists. For instance, if an AI Ops agent is part of the system, an alert could trigger that agent to attempt an automatic remediation (like retrying a job or rolling back to last good data). Or it could simply notify the AI monitoring agent to focus attention on the issue. This is an emerging area – in practice we’d start with Slack/email, and later could add an automated agent that listens for alerts and takes predefined actions (with safety checks).

**Alert Policy Examples:** To make this concrete, consider a few policy entries: - **Policy 1: Transaction Sync Delay** – *Trigger:* No new transactions loaded by 6am (expected daily load time). *Notification:* Slack alert to #pipeline-alerts; if not resolved by 8am, email to Data Engineering team distribution. *Escalation:* If still unresolved by 10am, page the on-call engineer (critical, as it starts affecting users). - **Policy 2: Data Quality Anomaly** – *Trigger:* Any Great Expectations test failure in production (e.g. schema mismatch or >5% transactions uncategorized unexpectedly). *Notification:* Slack alert tagging @data-team with test details and a link to the validation report. *Escalation:* If the same issue occurs 3 days in a row (recurring problem), create a Jira ticket and email data team lead for deeper investigation (to ensure it’s being worked on). - **Policy 3: Forecast Accuracy Alert** – *Trigger:* Monthly forecast error > 15%. *Notification:* Slack message in #ai-model-monitoring with the metric and note to retrain if persistent. *Escalation:* If two consecutive months breach the threshold, escalate to both data science and product owner via email (since it may impact user experience). - **Policy 4: ML Drift Alert** – *Trigger:* Detected category distribution drift with p-value < 0.001 (significant drift). *Notification:* Slack alert with a summary (which feature drifted) and link to an Evidently drift report for details. *Escalation:* If drift continues in next cycle, schedule a review meeting (could be manual process outside system) to plan model update.

All alerts should be aggregated into a central **incident management dashboard** as well, where their status (acknowledged, resolved) can be tracked. Grafana or other monitoring tools can list active alerts. It's good to also keep an **alert log** or report: e.g., a weekly report of how many alerts fired, to identify noisy ones or recurring issues so the team can improve the system (perhaps adjusting thresholds or fixing root causes).

Finally, when alerts fire, ensure there are **playbooks** available – i.e., runbooks on how to handle them. For example, if an alert says “Orphan payout detected,” the on-call team should have a document that guides them to check the specific table for null foreign keys, etc., so they can quickly mitigate (maybe by re-running an import or informing the user of a temporary glitch). The combination of proactive alerting and clear processes to address alerts closes the loop on the observability system.

---

### Sources:

1. Michael Segner, *“Data Validation Testing: Techniques, Examples, & Tools,”* Monte Carlo Data Blog – discusses data validation methods like range, type, uniqueness, and referential integrity checks <sup>1</sup> <sup>3</sup> <sup>2</sup> <sup>8</sup> .
2. Lasya, *“Monitoring & Alerting for Data Pipelines — Make Failures Visible, Fast,”* Medium (May 24, 2025) – outlines what to monitor in data pipelines (operational, data, system metrics) and tools integration like Airflow, Great Expectations, Slack alerts <sup>11</sup> <sup>38</sup> <sup>26</sup> .
3. Orchestra (Hugo Lu), *“Data Observability Open Source: Best Tools and Frameworks,”* (July 19, 2023) – recommends open-source observability tools (Prometheus, Grafana, OpenTelemetry) and steps to set up monitoring and alerting <sup>19</sup> .
4. Evidently AI, *“What is data drift in ML, and how to detect and handle it,”* (Jan 9, 2025) – explains data drift and methods to detect distribution changes using statistical tests and metrics <sup>23</sup> <sup>22</sup> .
5. Monte Carlo Data, *“The Ultimate Guide to Data Lineage,”* – describes the importance of lineage for tracing errors and compliance. Notably: *“When a dashboard shows incorrect numbers or a pipeline fails, lineage acts as your debugging roadmap”* <sup>36</sup> and that lineage provides audit trails for how metrics are calculated <sup>31</sup> .
6. Rubrik, *“AI Issues? Take Control with Rubrik Agent Rewind,”* (Press release, Aug 11, 2025) – highlights the need for auditing AI agent actions with an immutable trail from prompts to outcomes <sup>33</sup> .
7. Grafana OnCall Documentation, *“Escalation chains and routes for OnCall OSS,”* (2025) – details setting up alert routing and escalation policies (e.g., threshold-based and repeated escalation) in an open-source on-call tool <sup>43</sup> .

---

<sup>1</sup> <sup>2</sup> <sup>3</sup> <sup>4</sup> <sup>5</sup> <sup>6</sup> <sup>7</sup> <sup>8</sup> <sup>10</sup> <sup>40</sup> Data Validation Testing: Techniques, Examples, & Tools  
<https://www.montecarlodata.com/blog-data-validation-testing/>

9 A Pragmatic CDO's Field Guide to Data Quality — Part 7 — Finance Data Quality Playbook: Ledgers, Reconciliations, Controls, Model Risk | by Adnan Masood, PhD. | Sep, 2025 | Medium  
<https://medium.com/@adnanmasood/a-pragmatic-cdos-field-guide-to-data-quality-part-7-finance-data-quality-playbook-ledgers-2bb975683605>

11 12 13 14 15 16 25 26 28 37 38 39 Day 22: Monitoring & Alerting for Data Pipelines — Make Failures Visible, Fast | by Lasya | Medium  
<https://medium.com/@lasyachowdary1703/day-22-monitoring-alerting-for-data-pipelines-make-failures-visible-fast-b7863ae09e36>

17 18 19 27 Data Observability Open Source: Best Tools and Frameworks | Orchestra  
<https://www.getorchestra.io/guides/data-observability-open-source-best-tools-and-frameworks>

20 21 Model monitoring for ML in production: a comprehensive guide  
<https://www.evidentlyai.com/ml-in-production/model-monitoring>

22 23 What is data drift in ML, and how to detect and handle it  
<https://www.evidentlyai.com/ml-in-production/data-drift>

24 Detecting Data Drift: Tools, Techniques, and Best Practices for 2025  
<https://medium.com/@Symufolk/detecting-data-drift-tools-techniques-and-best-practices-for-2025-d0d34b863c3e>

29 30 31 32 36 The Ultimate Guide To Data Lineage  
<https://www.montecarlodata.com/blog-data-lineage/>

33 34 35 AI Issues? Take Control with Rubrik Agent Rewind | Rubrik  
<https://www.rubrik.com/insights/ai-issues-take-control-with-rubrik-agent-rewind>

41 42 43 Escalation chains and routes for OnCall OSS | Grafana OnCall documentation  
<https://grafana.com/docs/oncall/latest/configure/escalation-chains-and-routes/>